

# Replit Agent Implementation Instructions

## AI-Orchestrated Video Production Pipeline

---

### CRITICAL CONTEXT

You are implementing a video production system where **Claude (Anthropic API)** acts as the Project Lead. Claude doesn't just execute—it **plans, evaluates, and makes creative decisions**.

The current implementation has individual features that don't coordinate. Your job is to build the orchestration layer.

---

### IMPLEMENTATION PRIORITY ORDER

#### Phase 1: Scene Manifest System (Do This First)

The entire system hinges on Claude creating a structured **Scene Manifest** from the script. Without this, everything else is disconnected.

```
typescript
```

```
// /api/services/manifestGenerator.ts

import Anthropic from '@anthropic-ai/sdk';

const MANIFEST_SYSTEM_PROMPT = `You are an AI Video Producer. When given a marketing script, you create a detailed manifest for it. Your manifest must be executable—every visual requirement must include:
```

Your manifest must be executable—every visual requirement must include:

1. A detailed AI generation prompt (for Runway/Stability)
2. Search terms for stock footage fallback
3. Quality criteria to evaluate the result

Think like a professional video producer making creative decisions.';

```
export async function generateManifest(
  script: string,
  config: {
    platform: 'youtube' | 'tiktok' | 'instagram' | 'linkedin';
    duration: number;
    style: 'professional' | 'casual' | 'cinematic' | 'medical';
    targetAudience?: string;
  }
): Promise<SceneManifest> {
  const client = new Anthropic({ apiKey: process.env.ANTHROPIC_API_KEY });

  const response = await client.messages.create({
    model: 'claude-sonnet-4-20250514',
    max_tokens: 8000,
    system: MANIFEST_SYSTEM_PROMPT,
    messages: [
      {
        role: 'user',
        content: 'Create a Scene Manifest for this video:
```

SCRIPT:

`\${script}

CONFIGURATION:

- Platform: \${config.platform}
- Duration: \${config.duration} seconds
- Style: \${config.style}
- Audience: \${config.targetAudience || 'General'}

Return ONLY valid JSON matching this structure:

{

```
"videoId": "unique-id",
"metadata": { "title": "", "duration": 0, "platform": "", "aspectRatio": "", "style": "" },
"scenes": [
{
  "id": "scene-1",
  "section": "HOOK",
  "script": "The exact script text for this section",
  "voiceover": {
    "text": "Voiceover text",
    "estimatedDuration": 5,
    "emotion": "intrigued",
    "pacing": "normal"
  },
  "visualRequirements": {
    "primary": {
      "type": "ai-generated",
      "description": "Detailed description",
      "aiPrompt": "Exact prompt for Runway/Stability",
      "searchTerms": ["fallback", "search", "terms"],
      "duration": 5,
      "priority": "critical",
      "fallbackStrategy": "search stock footage for..."
    },
    "bRoll": [],
    "overlays": []
  },
  "transitions": { "in": "fade", "out": "crossfade" }
},
"qualityCriteria": {
  "mustHave": ["professional appearance", "matches brand"],
  "niceToHave": ["motion", "engaging composition"]
}
},
],
"globalAssets": {
  "backgroundMusic": { "mood": "uplifting", "intensity": "moderate" },
  "colorGrading": "natural",
  "typography": { "font": "clean sans-serif", "style": "minimal" }
}
},
});
```

```
    return JSON.parse(response.content[0].text);
}
```

## Phase 2: Asset Generation with Fallbacks

Each asset type needs a primary generator and fallback chain:

```
typescript
```

```
// /api/services/assetGenerator.ts

export class AssetGenerator {
    private runway: RunwayClient;
    private stability: StabilityClient;
    private huggingface: HuggingFaceClient;
    private pexels: PexelsClient;
    private elevenlabs: ElevenLabsClient;

    constructor() {
        this.runway = new RunwayClient(process.env.RUNWAY_API_KEY);
        this.stability = new StabilityClient(process.env.STABILITY_API_KEY);
        this.huggingface = new HuggingFaceClient(process.env.HUGGINGFACE_API_TOKEN);
        this.pexels = new PexelsClient(process.env.PEXELS_API_KEY);
        this.elevenlabs = new ElevenLabsClient(process.env.ELEVENLABS_API_KEY);
    }

    async generateVisual(requirement: VisualRequirement): Promise<GeneratedAsset> {
        const errors: Error[] = [];

        // Try primary: AI generation
        if (requirement.type === 'ai-generated' && requirement.aiPrompt) {
            try {
                // For longer durations, use Runway video
                if (requirement.duration > 3) {
                    return await this.generateRunwayVideo(requirement);
                }
                // For stills, use Stability
                return await this.generateStabilityImage(requirement);
            } catch (e) {
                errors.push(e);
                console.log(`Primary generation failed: ${e.message}`);
            }
        }

        // Fallback 1: HuggingFace
        try {
            return await this.generateHuggingFaceImage(requirement);
        } catch (e) {
            errors.push(e);
            console.log(`HuggingFace fallback failed: ${e.message}`);
        }
    }
}
```

```
// Fallback 2: Stock footage
try {
    return await this.searchStock(requirement);
} catch (e) {
    errors.push(e);
    console.log(`Stock fallback failed: ${e.message}`);
}

throw new Error(`All generation methods failed: ${errors.map(e => e.message).join(', ')}`);
}

private async generateRunwayVideo(req: VisualRequirement): Promise<GeneratedAsset> {
    // Runway Gen-4 Turbo API
    const task = await fetch('https://api.dev.runwayml.com/v1/image_to_video', {
        method: 'POST',
        headers: {
            'Authorization': `Bearer ${process.env.RUNWAY_API_KEY}`,
            'Content-Type': 'application/json',
            'X-Runway-Version': '2024-11-06'
        },
        body: JSON.stringify({
            model: 'gen4_turbo',
            promptText: req.aiPrompt,
            duration: Math.min(req.duration, 10),
            ratio: '16:9' // or '9:16' for vertical
        })
    });

    const { id } = await response.json();

    // Poll for completion
    let result;
    do {
        await sleep(5000);
        const pollRes = await fetch(`https://api.dev.runwayml.com/v1/tasks/${id}`, {
            headers: { 'Authorization': `Bearer ${process.env.RUNWAY_API_KEY}` }
        });
        result = await pollRes.json();
    } while (result.status === 'RUNNING' || result.status === 'PENDING');

    if (result.status !== 'SUCCEEDED') {
        throw new Error(`Runway failed: ${result.failure}`);
    }
}
```

```
        return {
          id: id,
          type: 'video',
          url: result.output[0],
          source: 'runway',
          metadata: { prompt: req.aiPrompt, duration: req.duration }
        };
      }

private async generateStabilityImage(req: VisualRequirement): Promise<GeneratedAsset> {
  const response = await fetch('https://api.stability.ai/v2beta/stable-image/generate/sd3', {
    method: 'POST',
    headers: {
      'Authorization': `Bearer ${process.env.STABILITY_API_KEY}`,
      'Content-Type': 'application/json',
      'Accept': 'application/json'
    },
    body: JSON.stringify({
      prompt: req.aiPrompt,
      negative_prompt: 'blurry, low quality, distorted, amateur',
      aspect_ratio: '16:9',
      output_format: 'png'
    })
  });
}

const data = await response.json();

return {
  id: crypto.randomUUID(),
  type: 'image',
  base64: data.image,
  source: 'stability',
  metadata: { prompt: req.aiPrompt }
};

async generateVoiceover(text: string, config: VoiceConfig): Promise<AudioAsset> {
  const response = await fetch(`https://api.elevenlabs.io/v1/text-to-speech/${config.voiceId || 'pNInz6obpgDQGcFmaJgB'}`)
  method: 'POST',
  headers: {
    'xi-api-key': process.env.ELEVENLABS_API_KEY,
    'Content-Type': 'application/json'
  },
  body: JSON.stringify({
```

```
text: text,
model_id: 'eleven_multilingual_v2',
voice_settings: {
  stability: 0.5,
  similarity_boost: 0.75,
  style: config.emotion === 'excited' ? 0.5 : 0.2
}
})
});
```

```
const buffer = await response.arrayBuffer();
```

```
return {
  id: crypto.randomUUID(),
  type: 'audio',
  buffer: Buffer.from(buffer),
  source: 'elevenlabs',
  metadata: { text }
};
```

### Phase 3: Quality Evaluation Gate

**Every generated asset must pass Claude's evaluation before use:**

typescript

```

// /api/services/qualityEvaluator.ts

export async function evaluateAsset(
  asset: GeneratedAsset,
  requirement: VisualRequirement,
  sceneContext: string
): Promise<QualityEvaluation> {
  const client = new Anthropic({ apiKey: process.env.ANTHROPIC_API_KEY });

  const content: MessageContent[] = [];

  // Include the image if we have it
  if (asset.base64) {
    content.push({
      type: 'image',
      source: { type: 'base64', media_type: 'image/png', data: asset.base64 }
    });
  }

  content.push({
    type: 'text',
    text: 'Evaluate this generated asset for a marketing video.

```

REQUIREMENT:

`${JSON.stringify(requirement, null, 2)}`

SCENE CONTEXT:

`${sceneContext}`

Score each dimension 0-100:

- Relevance: Does it match what was requested?
- Technical Quality: Resolution, composition, no artifacts?
- Brand Alignment: Professional, matches intended style?
- Emotional Impact: Does it support the marketing message?

Assets need 70+ overall score to be approved.

Return JSON:

```
{
  "scores": { "relevance": 0, "technicalQuality": 0, "brandAlignment": 0, "emotionalImpact": 0 },
  "overall": 0,
  "decision": "accept" | "regenerate" | "use-fallback",
  "reasoning": "Explain your decision",
}
```

```
"improvements": ["If regenerating, specific improvements needed"]  
}  
});  
  
const response = await client.messages.create({  
  model: 'claude-sonnet-4-20250514',  
  max_tokens: 1000,  
  messages: [{ role: 'user', content }]  
});  
  
return JSON.parse(response.content[0].text);  
}
```

## Phase 4: Production Orchestrator

The main class that ties everything together:

```
typescript
```

```
// /api/services/videoProducer.ts

import { EventEmitter } from 'events';

export class VideoProducer extends EventEmitter {
  private manifestGenerator: ManifestGenerator;
  private assetGenerator: AssetGenerator;
  private qualityEvaluator: QualityEvaluator;

  async produceVideo(script: string, config: VideoConfig): Promise<ProductionResult> {
    const productionId = crypto.randomUUID();

    // PHASE 1: Analyze Script
    this.emit('phase', { phase: 'analysis', status: 'started' });
    this.emit('log', { type: 'decision', message: 'AI Producer analyzing script...' });

    const manifest = await this.manifestGenerator.generate(script, config);

    this.emit('log', {
      type: 'decision',
      message: `Created ${manifest.scenes.length} scenes with detailed asset requirements`
    });
    this.emit('phase', { phase: 'analysis', status: 'complete', manifest });

    // PHASE 2: Generate All Assets
    this.emit('phase', { phase: 'generation', status: 'started' });

    const generatedAssets: GeneratedAsset[] = [];

    // Generate voiceovers first (needed for timing)
    for (const scene of manifest.scenes) {
      this.emit('log', {
        type: 'generation',
        message: `Generating voiceover for ${scene.section}...`
      });

      const voiceover = await this.assetGenerator.generateVoiceover(
        scene.voiceover.text,
        { emotion: scene.voiceover.emotion }
      );

      generatedAssets.push({ ...voiceover, sceneId: scene.id, assetType: 'voiceover' });
    }
  }
}
```

```

// Generate visuals in parallel
const visualPromises = manifest.scenes.flatMap(scene => {
  const promises = [];

  // Primary visual
  promises.push(
    this.assetGenerator.generateVisual(scene.visualRequirements.primary)
      .then(asset => ({ ...asset, sceneId: scene.id, assetType: 'primary' }))
  );

  // B-roll
  scene.visualRequirements.bRoll.forEach((req, i) => {
    promises.push(
      this.assetGenerator.generateVisual(req)
        .then(asset => ({ ...asset, sceneId: scene.id, assetType: `broll-${i}` }))
    );
  });
}

return promises;
});

const visualResults = await Promise.allSettled(visualPromises);

for (const result of visualResults) {
  if (result.status === 'fulfilled') {
    generatedAssets.push(result.value);
    this.emit('log', {
      type: 'generation',
      message: `Generated ${result.value.assetType} for scene ${result.value.sceneId}`
    });
  } else {
    this.emit('log', {
      type: 'error',
      message: `Generation failed: ${result.reason.message}`
    });
  }
}

this.emit('phase', { phase: 'generation', status: 'complete' });

// PHASE 3: Quality Evaluation
this.emit('phase', { phase: 'evaluation', status: 'started' });

```

```

const approved: GeneratedAsset[] = [];
const needsRegeneration: { asset: GeneratedAsset; evaluation: QualityEvaluation }[] = [];

for (const asset of generatedAssets.filter(a => a.type !== 'audio')) {
  const scene = manifest.scenes.find(s => s.id === asset.sceneId);
  const requirement = this.getRequirement(scene, asset.assetType);

  this.emit('log', {
    type: 'evaluation',
    message: `Evaluating ${asset.assetType} for ${scene.section}...`
  });

  const evaluation = await this.qualityEvaluator.evaluate(
    asset,
    requirement,
    scene.script
  );

  if (evaluation.decision === 'accept') {
    approved.push({ ...asset, evaluation });
    this.emit('log', {
      type: 'success',
      message: `✓ ${scene.section} ${asset.assetType}: ${evaluation.overall}/100 - Approved`
    });
  } else {
    needsRegeneration.push({ asset, evaluation });
    this.emit('log', {
      type: 'evaluation',
      message: `⚠ ${scene.section} ${asset.assetType}: ${evaluation.overall}/100 - ${evaluation.reasoning}`
    });
  }
}

// Add all audio assets (not evaluated visually)
approved.push(...generatedAssets.filter(a => a.type === 'audio'));

this.emit('phase', { phase: 'evaluation', status: 'complete' });

// PHASE 4: Handle Failures (Regenerate or Fallback)
if (needsRegeneration.length > 0) {
  this.emit('phase', { phase: 'regeneration', status: 'started' });

  for (const { asset, evaluation } of needsRegeneration) {
    const scene = manifest.scenes.find(s => s.id === asset.sceneId);
  }
}

```

```

const requirement = this.getRequirement(scene, asset.assetType);

this.emit('log', {
  type: 'generation',
  message: `Regenerating ${asset.assetType} with improvements: ${evaluation.improvements?.join(', ')}`});
});

// Modify the prompt based on evaluation feedback
const improvedRequirement = {
  ...requirement,
  aiPrompt: `${requirement.aiPrompt}. IMPORTANT: ${evaluation.improvements?.join('. ')}`;
};

try {
  const regenerated = await this.assetGenerator.generateVisual(improvedRequirement);
  const reEvaluation = await this.qualityEvaluator.evaluate(regenerated, requirement, scene.script);

  if (reEvaluation.decision === 'accept') {
    approved.push({ ...regenerated, scenId: asset.scenId, assetType: asset.assetType, evaluation: reEvaluation });
    this.emit('log', { type: 'success', message: `✓ Regenerated asset approved: ${reEvaluation.overall}/100` });
  } else {
    // Use fallback
    this.emit('log', { type: 'generation', message: `Using stock fallback for ${asset.assetType}...` });
    const fallback = await this.assetGenerator.searchStock(requirement);
    approved.push({ ...fallback, scenId: asset.scenId, assetType: asset.assetType });
  }
} catch (e) {
  this.emit('log', { type: 'error', message: `Regeneration failed, using placeholder` });
}

this.emit('phase', { phase: 'regeneration', status: 'complete' });
}

// PHASE 5: Assembly
this.emit('phase', { phase: 'assembly', status: 'started' });
this.emit('log', { type: 'success', message: 'All assets approved. Assembling final video...' });

const assembled = await this.assembleVideo(approved, manifest);

this.emit('phase', { phase: 'assembly', status: 'complete', video: assembled });
this.emit('log', { type: 'success', message: `Video complete! Duration: ${assembled.duration}s` });

return {

```

```
success: true,  
videoUrl: assembled.url,  
manifest,  
assets: approved  
};  
}  
}
```

## Phase 5: Real-Time UI Updates

Use Server-Sent Events or WebSocket to stream production progress:

```
typescript
```

```
// /api/routes/production.ts

import { Router } from 'express';
import { VideoProducer } from '../services/videoProducer';

const router = Router();

router.post('/produce', async (req, res) => {
  const { script, config } = req.body;

  // Set up SSE
  res.setHeader('Content-Type', 'text/event-stream');
  res.setHeader('Cache-Control', 'no-cache');
  res.setHeader('Connection', 'keep-alive');

  const producer = new VideoProducer();

  // Stream all events to client
  producer.on('phase', (data) => {
    res.write(`event: phase\ndata: ${JSON.stringify(data)}\n\n`);
  });

  producer.on('log', (data) => {
    res.write(`event: log\ndata: ${JSON.stringify({ ...data, timestamp: new Date() })}\n\n`);
  });

  try {
    const result = await producer.produceVideo(script, config);
    res.write(`event: complete\ndata: ${JSON.stringify(result)}\n\n`);

  } catch (error) {
    res.write(`event: error\ndata: ${JSON.stringify({ message: error.message })}\n\n`);
  }

  res.end();
});

export default router;
```

## Phase 6: Frontend Integration

tsx

```
// /components/VideoProducer.tsx

export function VideoProducer({ script, config }: Props) {
  const [phases, setPhases] = useState<Phase[]>([]);
  const [logs, setLogs] = useState<LogEntry[]>([]);
  const [result, setResult] = useState<ProductionResult | null>(null);

  const startProduction = useCallback(() => {
    const eventSource = new EventSource('/api/produce?script=${encodeURIComponent(script)}&config=${encodeURIComponent(config)}');

    eventSource.addEventListener('phase', (e) => {
      const data = JSON.parse(e.data);
      setPhases(prev => {
        const existing = prev.findIndex(p => p.phase === data.phase);
        if (existing >= 0) {
          const updated = [...prev];
          updated[existing] = data;
          return updated;
        }
        return [...prev, data];
      });
    });

    eventSource.addEventListener('log', (e) => {
      const data = JSON.parse(e.data);
      setLogs(prev => [...prev, data]);
    });

    eventSource.addEventListener('complete', (e) => {
      const data = JSON.parse(e.data);
      setResult(data);
      eventSource.close();
    });

    eventSource.addEventListener('error', () => {
      eventSource.close();
    });
  }, [script, config]);

  return (
    <div className="producer-container">
      <WorkflowProgress phases={phases} />
      <ProductionLog logs={logs} />
    </div>
  );
}
```

```
{result && <VideoPreview video={result} />
  <button onClick={startProduction}>Start Production</button>
</div>
);
}
```

## KEY IMPLEMENTATION RULES

### 1. Claude Is Always The Decision Maker

- Never generate assets without Claude first creating a manifest
- Never use assets without Claude evaluating them
- Let Claude modify prompts based on evaluation feedback

### 2. Always Have Fallbacks

```
Primary: Runway/Stability AI
↓ failure
Fallback 1: HuggingFace
↓ failure
Fallback 2: Stock footage (Pexels/Pixabay)
↓ failure
Fallback 3: Static placeholder with text
```

### 3. Stream Everything to UI

- Users must see every decision the AI makes
- Show asset previews as they're generated
- Allow users to override AI decisions

### 4. Quality Gates Are Non-Negotiable

- Every visual asset needs 70+ score
- Failed assets trigger regeneration with improvements
- Track why assets fail to improve prompts

### 5. Scene Manifest Is Sacred

- The manifest defines the entire video structure
- All assets reference their scene ID
- Assembly follows the manifest timeline exactly

---

## ENVIRONMENT VARIABLES REQUIRED

```
env

ANTHROPIC_API_KEY=      # Claude for orchestration
RUNWAY_API_KEY=         # Video generation
STABILITY_API_KEY=      # Image generation
ELEVENLABS_API_KEY=     # Voiceover
HUGGINGFACE_API_TOKEN=   # Fallback image generation
PEXELS_API_KEY=          # Stock video
PIXABAY_API_KEY=         # Stock video/images
UNSPLASH_ACCESS_KEY=    # Stock images
```

---

## TESTING CHECKLIST

- Manifest generates correctly from Pine Hill Farm script
  - All 5 scenes have visual requirements with AI prompts
  - Voiceovers generate with correct emotion/pacing
  - At least one AI-generated image passes quality gate
  - Failed assets trigger regeneration
  - Fallback to stock footage works
  - Real-time logs appear in UI
  - Final video assembles with all approved assets
- 

## COMMON ISSUES & FIXES

**Issue:** Runway returns 429 (rate limited) **Fix:** Implement exponential backoff, use stock footage as immediate fallback

**Issue:** Generated images don't match script **Fix:** Make Claude's AI prompts more specific, include scene context

**Issue:** Quality evaluation too strict **Fix:** Lower threshold to 60, or have Claude explain borderline cases

**Issue:** Assembly fails with missing assets **Fix:** Validate all required assets exist before assembly phase

---

This document gives Replit agents everything needed to build a properly orchestrated video production system.